

Lab10. Patients Physical Activities Prediction using Boosting

```
In [1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
from sklearn.metrics import accuracy_score, classification_report
from sklearn.ensemble import GradientBoostingClassifier, AdaBoostClassifier, RandomForestClassifier
from sklearn.linear_model import LogisticRegressionCV
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder
```

Step 1 [Understand Data]

```
In [2]: original = pd.read_csv("Human_Activity_Data.csv")
```

```
In [3]: original.head()
```

```
Out[3]:
```

	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X	tBodyAcc-mad()-Y	tBodyAcc-mad()-Z
0	0.288585	-0.020294	-0.132905	-0.995279	-0.983111	-0.913526	-0.995112	-0.983185	-0.983185
1	0.278419	-0.016411	-0.123520	-0.998245	-0.975300	-0.960322	-0.998807	-0.974914	-0.974914
2	0.279653	-0.019467	-0.113462	-0.995380	-0.967187	-0.978944	-0.996520	-0.963668	-0.963668
3	0.279174	-0.026201	-0.123283	-0.996091	-0.983403	-0.990675	-0.997099	-0.982750	-0.982750
4	0.276629	-0.016570	-0.115362	-0.998139	-0.980817	-0.990482	-0.998321	-0.979672	-0.979672

5 rows × 562 columns

```
In [4]: original.columns
```

```
Out[4]: Index(['tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z',
               'tBodyAcc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z',
               'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z',
               'tBodyAcc-max()-X',
               ...,
               'fBodyBodyGyroJerkMag-skewness()', 'fBodyBodyGyroJerkMag-kurtosis()',
               'angle(tBodyAccMean,gravity)', 'angle(tBodyAccJerkMean,gravityMean)',
               'angle(tBodyGyroMean,gravityMean)',
               'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)',
               'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'Activity'],
              dtype='object', length=562)
```

In [5]: `original.shape`

Out[5]: (10299, 562)

In [6]: `original.dtypes`

```
Out[6]: tBodyAcc-mean()-X      float64
tBodyAcc-mean()-Y      float64
tBodyAcc-mean()-Z      float64
tBodyAcc-std()-X      float64
tBodyAcc-std()-Y      float64
...
angle(tBodyGyroJerkMean,gravityMean) float64
angle(X,gravityMean)      float64
angle(Y,gravityMean)      float64
angle(Z,gravityMean)      float64
Activity                  object
Length: 562, dtype: object
```

In [7]: `original.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10299 entries, 0 to 10298
Columns: 562 entries, tBodyAcc-mean()-X to Activity
dtypes: float64(561), object(1)
memory usage: 44.2+ MB
```

In [8]: `original.value_counts()`

```
Out[8]: tBodyAcc-mean()-X tBodyAcc-mean()-Y tBodyAcc-mean()-Z tBodyAcc-std()-X tB
odyAcc-std()-Y tBodyAcc-std()-Z tBodyAcc-mad()-X tBodyAcc-mad()-Y tBodyAc
c-mad()-Z tBodyAcc-max()-X tBodyAcc-max()-Y tBodyAcc-max()-Z tBodyAcc-min
()-X tBodyAcc-min()-Y tBodyAcc-min()-Z tBodyAcc-sma() tBodyAcc-energy()-X
tBodyAcc-energy()-Y tBodyAcc-energy()-Z tBodyAcc-iqr()-X tBodyAcc-iqr()-Y
tBodyAcc-iqr()-Z tBodyAcc-entropy()-X tBodyAcc-entropy()-Y tBodyAcc-entrop
y()-Z tBodyAcc-arCoeff()-X,1 tBodyAcc-arCoeff()-X,2 tBodyAcc-arCoeff()-X,3
tBodyAcc-arCoeff()-X,4 tBodyAcc-arCoeff()-Y,1 tBodyAcc-arCoeff()-Y,2 tBody
Acc-arCoeff()-Y,3 tBodyAcc-arCoeff()-Y,4 tBodyAcc-arCoeff()-Z,1 tBodyAcc-a
rCoeff()-Z,2 tBodyAcc-arCoeff()-Z,3 tBodyAcc-arCoeff()-Z,4 tBodyAcc-correl
ation()-X,Y tBodyAcc-correlation()-X,Z tBodyAcc-correlation()-Y,Z tGravity
Acc-mean()-X tGravityAcc-mean()-Y tGravityAcc-mean()-Z tGravityAcc-std()-X
tGravityAcc-std()-Y tGravityAcc-std()-Z tGravityAcc-mad()-X tGravityAcc-ma
d()-Y tGravityAcc-mad()-Z tGravityAcc-max()-X tGravityAcc-max()-Y tGravit
yAcc-max()-Z tGravityAcc-min()-X tGravityAcc-min()-Y tGravityAcc-min()-Z
tGravityAcc-sma() tGravityAcc-energy()-X tGravityAcc-energy()-Y tGravityAc
c-energy()-Z tGravityAcc-iqr()-X tGravityAcc-iqr()-Y tGravityAcc-iqr()-Z
tGravityAcc-entropy()-X tGravityAcc-entropy()-Y tGravityAcc-entropy()-Z tG
ravityAcc-arCoeff()-X,1 tGravityAcc-arCoeff()-X,2 tGravityAcc-arCoeff()-X,3
```

In [9]: `label_encoder = LabelEncoder()`
`original["label_Activity"] = label_encoder.fit_transform(original["Activity"])`

Step2. [Build a small dataset]

```
In [10]: original.Activity.value_counts()
```

```
Out[10]: LAYING                1944
          STANDING            1906
          SITTING             1777
          WALKING             1722
          WALKING_UPSTAIRS    1544
          WALKING_DOWNSTAIRS  1406
          Name: Activity, dtype: int64
```

```
In [11]: original.label_Activity.value_counts()
```

```
Out[11]: 0    1944
          2    1906
          1    1777
          3    1722
          5    1544
          4    1406
          Name: label_Activity, dtype: int64
```

Take first 3000 samples for each 6 activities and build classifier

```
In [12]: tem1 = original[original['Activity']=='LAYING'][:500]
          tem2 = original[original['Activity']=='SITTING'][:500]
          tem3 = original[original['Activity']=='WALKING'][:500]
          tem4 = original[original['Activity']=='STANDING'][:500]
          tem5 = original[original['Activity']=='WALKING_UPSTAIRS'][:500]
          tem6 = original[original['Activity']=='WALKING_DOWNSTAIRS'][:500]
```

```
In [13]: new_df = pd.concat([tem1,tem2,tem3,tem4,tem5,tem6])
```

```
In [14]: new_df.to_csv("human_activity_clipped3000.csv")
```

```
In [15]: new_df = pd.read_csv("human_activity_clipped3000.csv")
```

```
In [16]: new_df.head()
```

```
Out[16]:
```

	Unnamed: 0	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X	tB
0	51	0.403474	-0.015074	-0.118167	-0.914811	-0.895231	-0.891748	-0.917696	-1
1	52	0.278373	-0.020561	-0.096825	-0.984883	-0.991118	-0.982112	-0.987985	-1
2	53	0.276555	-0.017869	-0.107621	-0.994195	-0.996372	-0.995615	-0.994901	-1
3	54	0.279575	-0.017276	-0.109481	-0.996135	-0.995812	-0.998689	-0.996393	-1
4	55	0.276527	-0.016819	-0.107983	-0.996775	-0.997256	-0.995422	-0.997167	-1

5 rows × 564 columns

```
In [17]: new_df.shape
```

```
Out[17]: (3000, 564)
```

```
In [18]: new_df.columns
```

```
Out[18]: Index(['Unnamed: 0', 'tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y',
               'tBodyAcc-mean()-Z', 'tBodyAcc-std()-X', 'tBodyAcc-std()-Y',
               'tBodyAcc-std()-Z', 'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y',
               'tBodyAcc-mad()-Z',
               ...,
               'fBodyBodyGyroJerkMag-kurtosis()', 'angle(tBodyAccMean,gravity)',
               'angle(tBodyAccJerkMean,gravityMean)',
               'angle(tBodyGyroMean,gravityMean)',
               'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)',
               'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'Activity',
               'label_Activity'],
              dtype='object', length=564)
```

```
In [19]: new_df.dtypes
```

```
Out[19]: Unnamed: 0          int64
         tBodyAcc-mean()-X    float64
         tBodyAcc-mean()-Y    float64
         tBodyAcc-mean()-Z    float64
         tBodyAcc-std()-X     float64
         ...
         angle(X,gravityMean)  float64
         angle(Y,gravityMean)  float64
         angle(Z,gravityMean)  float64
         Activity              object
         label_Activity        int64
         Length: 564, dtype: object
```

In [20]: `new_df.value_counts()`

Out[20]: Unnamed: 0 tBodyAcc-mean()-X tBodyAcc-mean()-Y tBodyAcc-mean()-Z tBodyAcc-std()-X tBodyAcc-std()-Y tBodyAcc-std()-Z tBodyAcc-mad()-X tBodyAcc-mad()-Y tBodyAcc-mad()-Z tBodyAcc-max()-X tBodyAcc-max()-Y tBodyAcc-max()-Z tBodyAcc-min()-X tBodyAcc-min()-Y tBodyAcc-min()-Z tBodyAcc-sma() tBodyAcc-energy()-X tBodyAcc-energy()-Y tBodyAcc-energy()-Z tBodyAcc-iqr()-X tBodyAcc-iqr()-Y tBodyAcc-iqr()-Z tBodyAcc-entropy()-X tBodyAcc-entropy()-Y tBodyAcc-entropy()-Z tBodyAcc-arCoeff()-X,1 tBodyAcc-arCoeff()-X,2 tBodyAcc-arCoeff()-X,3 tBodyAcc-arCoeff()-X,4 tBodyAcc-arCoeff()-Y,1 tBodyAcc-arCoeff()-Y,2 tBodyAcc-arCoeff()-Y,3 tBodyAcc-arCoeff()-Y,4 tBodyAcc-arCoeff()-Z,1 tBodyAcc-arCoeff()-Z,2 tBodyAcc-arCoeff()-Z,3 tBodyAcc-arCoeff()-Z,4 tBodyAcc-correlation()-X,Y tBodyAcc-correlation()-X,Z tBodyAcc-correlation()-Y,Z tGravityAcc-mean()-X tGravityAcc-mean()-Y tGravityAcc-mean()-Z tGravityAcc-std()-X tGravityAcc-std()-Y tGravityAcc-std()-Z tGravityAcc-mad()-X tGravityAcc-mad()-Y tGravityAcc-mad()-Z tGravityAcc-max()-X tGravityAcc-max()-Y tGravityAcc-max()-Z tGravityAcc-min()-X tGravityAcc-min()-Y tGravityAcc-min()-Z tGravityAcc-sma() tGravityAcc-energy()-X tGravityAcc-energy()-Y tGravityAcc-energy()-Z tGravityAcc-iqr()-X tGravityAcc-iqr()-Y tGravityAcc-iqr()-Z tGravityAcc-entropy()-X tGravityAcc-entropy()-Y tGravityAcc-entropy()-Z tGravityAcc-arCoeff()-X,1 tGravityAcc-arCoeff()-X,2 tGravityAcc-arCoeff()-X,3 tGravityAcc-arCoeff()-X,4 tGravityAcc-arCoeff()-Y,1 tGravityAcc-arCoeff()-Y,2 tGravityAcc-arCoeff()-Y,3 tGravityAcc-arCoeff()-Y,4 tGravityAcc-arCoeff()-Z,1 tGravityAcc-arCoeff()-Z,2 tGravityAcc-arCoeff()-Z,3 tGravityAcc-arCoeff()-Z,4

In [21]: `X=new_df.drop(['Activity','label_Activity'],axis=1)`
`y=new_df.label_Activity`
`X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_s`

Build GradientBoostingClassifier for 3000 samples

In [22]: `model_ = GradientBoostingClassifier(subsample=0.5,n_estimators=100,learning_rate=`
`model_.fit(X_train,y_train)`
`y_pred=model_.predict(X_test)`

In [23]: `print(accuracy_score(y_test,y_pred))`
`print(classification_report(y_test,y_pred))`

0.705

	precision	recall	f1-score	support
0	0.36	0.11	0.16	94
1	0.57	0.70	0.63	97
2	0.54	0.74	0.63	101
3	0.82	0.92	0.87	105
4	0.87	0.87	0.87	103
5	0.88	0.83	0.86	100
accuracy			0.70	600
macro avg	0.68	0.70	0.67	600
weighted avg	0.68	0.70	0.68	600

Build AdaBoostClassifier for 3000 samples

```
In [24]: base = DecisionTreeClassifier(max_features=4)
base2 = AdaBoostClassifier(base_estimator=base, random_state=0)
param_grid = {'n_estimators': [100, 150, 200], 'learning_rate': [0.01, 0.001]}
```

```
In [25]: model2_ = GridSearchCV(base2, param_grid, cv=10, n_jobs=-1)
model2_.fit(X__train, y__train)
y__pred2=model2_.predict(X__test)
```

```
In [26]: print(accuracy_score(y__test, y__pred2))
print(classification_report(y__test, y__pred2))
```

```
0.7716666666666666
              precision    recall  f1-score   support

     0       0.76       0.69       0.72         94
     1       0.62       0.65       0.63         97
     2       0.69       0.71       0.70        101
     3       0.82       0.92       0.87        105
     4       0.90       0.81       0.85        103
     5       0.86       0.83       0.85        100

 accuracy                   0.77         600
 macro avg              0.77       0.77       0.77         600
 weighted avg           0.78       0.77       0.77         600
```

```
In [27]: model2_.best_estimator_
```

```
Out[27]: AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_features=4),
                             learning_rate=0.01, n_estimators=100, random_state=0)
```

Build LogisticRegressionCV classifier for 3000 samples

```
In [28]: model3_ = LogisticRegressionCV(cv=4, Cs=5, penalty='l2')
model3_.fit(X__train, y__train)
y__pred3=model3_.predict(X__test)
```

```
In [29]: print(accuracy_score(y__test,y__pred3))
print(classification_report(y__test,y__pred3))
```

```
0.9766666666666667
      precision    recall  f1-score   support

0         0.99        1.00        0.99         94
1         0.96        0.92        0.94         97
2         0.93        0.96        0.95        101
3         1.00        0.99        1.00        105
4         0.98        1.00        0.99        103
5         1.00        0.99        0.99        100

 accuracy          0.98        600
 macro avg         0.98        0.98        0.98        600
 weighted avg      0.98        0.98        0.98        600
```

Find Best no. of trees and Best Learning Rate using Grid Search and Cross Validation for 3000 samples

```
In [30]: Param_grid={'n_estimators':[50,100,200, 400], 'learning_rate':[0.1,0.01]}
```

```
In [31]: all_scores = cross_val_score(estimator=model_,X=X__train,y=y__train,cv=5)
print(all_scores)
```

```
[0.10416667 0.55416667 0.72708333 0.80416667 0.95      ]
```

```
In [32]: model4_ = GridSearchCV(estimator=model_,param_grid=Param_grid,cv=5,n_jobs=-1)
model4_.fit(X__train,y__train)
y__pred4=model4_.predict(X__test)
```

```
In [33]: print(accuracy_score(y__test,y__pred4))
print(classification_report(y__test,y__pred4))
```

```
0.985
      precision    recall  f1-score   support

0         1.00        1.00        1.00         94
1         0.97        0.95        0.96         97
2         0.95        0.97        0.96        101
3         1.00        0.99        1.00        105
4         0.99        1.00        1.00        103
5         1.00        1.00        1.00        100

 accuracy          0.98        600
 macro avg         0.99        0.98        0.98        600
 weighted avg      0.99        0.98        0.98        600
```

```
In [34]: model4_.best_estimator_
```

```
Out[34]: GradientBoostingClassifier(max_depth=1, max_features=4, n_estimators=400,
                                     random_state=0, subsample=0.5)
```

Build VotingClassifier for 3000 samples

```
In [35]: model5_ = VotingClassifier(estimators=[('lr', model3_), ('gbc', base2)], voting='soft')
model5_.fit(X__train, y__train)
y__pred5 = model5_.predict(X__test)
```

```
In [36]: print(accuracy_score(y__test, y__pred5))
print(classification_report(y__test, y__pred5))
```

```
0.7716666666666666
```

	precision	recall	f1-score	support
0	0.76	0.69	0.72	94
1	0.62	0.65	0.63	97
2	0.69	0.71	0.70	101
3	0.82	0.92	0.87	105
4	0.90	0.81	0.85	103
5	0.86	0.83	0.85	100
accuracy			0.77	600
macro avg	0.77	0.77	0.77	600
weighted avg	0.78	0.77	0.77	600

From this 3000 samples you should take 1500 samples for each activities

```
In [37]: temp1 = new_df[new_df['Activity']=='LAYING'][:500]
temp2 = new_df[new_df['Activity']=='SITTING'][:500]
temp3 = new_df[new_df['Activity']=='WALKING'][:500]
```

```
In [38]: df = pd.concat([temp1, temp2, temp3])
```

```
In [39]: df.to_csv("human_activity_clipped1500.csv")
```


```
In [40]: df = pd.read_csv("human_activity_clipped1500.csv")
```


In [41]: `df.head()`

Out[41]:

	Unnamed: 0	Unnamed: 0.1	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X	tBodyAcc-mad()-Y	tBodyAcc-mad()-Z
0	0	51	0.403474	-0.015074	-0.118167	-0.914811	-0.895231	-0.891748	-0.914811	-0.895231	-0.891748
1	1	52	0.278373	-0.020561	-0.096825	-0.984883	-0.991118	-0.982112	-0.984883	-0.991118	-0.982112
2	2	53	0.276555	-0.017869	-0.107621	-0.994195	-0.996372	-0.995615	-0.994195	-0.996372	-0.995615
3	3	54	0.279575	-0.017276	-0.109481	-0.996135	-0.995812	-0.998689	-0.996135	-0.995812	-0.998689
4	4	55	0.276527	-0.016819	-0.107983	-0.996775	-0.997256	-0.995422	-0.996775	-0.997256	-0.995422

5 rows × 565 columns



In [42]: `df.shape`

Out[42]: (1500, 565)

In [43]: `df.columns`

Out[43]: Index(['Unnamed: 0', 'Unnamed: 0.1', 'tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z', 'tBodyAcc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z', 'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z', 'tBodyGyro-mean()-X', 'tBodyGyro-mean()-Y', 'tBodyGyro-mean()-Z', 'tBodyGyro-std()-X', 'tBodyGyro-std()-Y', 'tBodyGyro-std()-Z', 'tBodyGyro-mad()-X', 'tBodyGyro-mad()-Y', 'tBodyGyro-mad()-Z', 'tBodyGyroJerk-mean()-X', 'tBodyGyroJerk-mean()-Y', 'tBodyGyroJerk-mean()-Z', 'tBodyGyroJerk-std()-X', 'tBodyGyroJerk-std()-Y', 'tBodyGyroJerk-std()-Z', 'tBodyGyroJerk-mad()-X', 'tBodyGyroJerk-mad()-Y', 'tBodyGyroJerk-mad()-Z', 'fBodyBodyGyroJerkMag-kurtosis()', 'angle(tBodyAccMean,gravity)', 'angle(tBodyAccJerkMean,gravityMean)', 'angle(tBodyGyroMean,gravityMean)', 'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)', 'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'Activity', 'label_Activity'], dtype='object', length=565)

In [44]: `df.dtypes`

Out[44]:

Unnamed: 0	int64
Unnamed: 0.1	int64
tBodyAcc-mean()-X	float64
tBodyAcc-mean()-Y	float64
tBodyAcc-mean()-Z	float64
...	...
angle(X,gravityMean)	float64
angle(Y,gravityMean)	float64
angle(Z,gravityMean)	float64
Activity	object
label_Activity	int64
Length: 565, dtype: object	

In [45]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Columns: 565 entries, Unnamed: 0 to label_Activity
dtypes: float64(561), int64(3), object(1)
memory usage: 6.5+ MB
```

In [46]: `df.value_counts()`

Out[46]: Unnamed: 0 Unnamed: 0.1 tBodyAcc-mean()-X tBodyAcc-mean()-Y tBodyAcc-mean()
()-Z tBodyAcc-std()-X tBodyAcc-std()-Y tBodyAcc-std()-Z tBodyAcc-mad()-X
tBodyAcc-mad()-Y tBodyAcc-mad()-Z tBodyAcc-max()-X tBodyAcc-max()-Y tBody
Acc-max()-Z tBodyAcc-min()-X tBodyAcc-min()-Y tBodyAcc-min()-Z tBodyAcc-s
ma() tBodyAcc-energy()-X tBodyAcc-energy()-Y tBodyAcc-energy()-Z tBodyAcc
-iqr()-X tBodyAcc-iqr()-Y tBodyAcc-iqr()-Z tBodyAcc-entropy()-X tBodyAcc-
entropy()-Y tBodyAcc-entropy()-Z tBodyAcc-arCoeff()-X,1 tBodyAcc-arCoeff()
-X,2 tBodyAcc-arCoeff()-X,3 tBodyAcc-arCoeff()-X,4 tBodyAcc-arCoeff()-Y,1
tBodyAcc-arCoeff()-Y,2 tBodyAcc-arCoeff()-Y,3 tBodyAcc-arCoeff()-Y,4 tBody
Acc-arCoeff()-Z,1 tBodyAcc-arCoeff()-Z,2 tBodyAcc-arCoeff()-Z,3 tBodyAcc-a
rCoeff()-Z,4 tBodyAcc-correlation()-X,Y tBodyAcc-correlation()-X,Z tBodyAc
c-correlation()-Y,Z tGravityAcc-mean()-X tGravityAcc-mean()-Y tGravityAcc-
mean()-Z tGravityAcc-std()-X tGravityAcc-std()-Y tGravityAcc-std()-Z tGra
vityAcc-mad()-X tGravityAcc-mad()-Y tGravityAcc-mad()-Z tGravityAcc-max()-
X tGravityAcc-max()-Y tGravityAcc-max()-Z tGravityAcc-min()-X tGravityAcc
-min()-Y tGravityAcc-min()-Z tGravityAcc-sma() tGravityAcc-energy()-X tGr
avityAcc-energy()-Y tGravityAcc-energy()-Z tGravityAcc-iqr()-X tGravityAcc
-iqr()-Y tGravityAcc-iqr()-Z tGravityAcc-entropy()-X tGravityAcc-entropy()
-Y tGravityAcc-entropy()-Z tGravityAcc-arCoeff()-X,1 tGravityAcc-arCoeff()

Step3. [Build GradientBoostingClassifier]

In [47]: `X=df.drop(['Activity','label_Activity'],axis=1)`
`y=df.label_Activity`

In [48]: `X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=2)`

In [49]: `model = GradientBoostingClassifier(subsample=0.5,n_estimators=100,learning_rate=1)`

In [50]: `model.fit(X_train,y_train)`

Out[50]: GradientBoostingClassifier(learning_rate=1.0, max_depth=1, max_features=4,
random_state=42, subsample=0.5)

```
In [51]: y_pred=model.predict(X_test)
y_pred
```

```
Out[51]: array([0, 1, 1, 1, 3, 1, 0, 0, 3, 3, 1, 3, 3, 3, 3, 1, 3, 3, 3, 0, 3, 1,
0, 1, 0, 0, 0, 1, 0, 1, 3, 0, 3, 1, 1, 3, 0, 0, 3, 1, 0, 1, 0, 0,
3, 0, 3, 3, 3, 1, 0, 0, 0, 1, 1, 0, 3, 3, 0, 3, 1, 1, 0, 1, 0, 3,
3, 0, 3, 0, 0, 3, 0, 0, 1, 3, 3, 0, 3, 0, 3, 3, 1, 3, 1, 1, 1, 3,
0, 3, 1, 3, 1, 1, 1, 3, 1, 0, 3, 1, 3, 0, 0, 3, 1, 3, 1, 0, 1, 0,
1, 0, 0, 0, 1, 1, 1, 1, 3, 0, 0, 0, 0, 0, 1, 3, 3, 1, 1, 1, 1, 3,
0, 1, 0, 3, 3, 3, 3, 3, 3, 1, 0, 3, 1, 1, 3, 3, 3, 0, 1, 1, 0, 0,
0, 3, 0, 1, 3, 3, 1, 1, 1, 1, 0, 1, 0, 0, 0, 3, 0, 1, 1, 3, 0, 0,
0, 0, 1, 3, 1, 0, 1, 1, 3, 1, 1, 0, 3, 0, 1, 3, 0, 3, 0, 1, 3, 0,
1, 1, 3, 1, 0, 0, 0, 3, 3, 1, 1, 3, 3, 3, 3, 1, 0, 1, 0, 0, 3, 0,
3, 3, 0, 0, 3, 1, 3, 0, 3, 1, 1, 3, 0, 1, 1, 1, 0, 0, 3, 1, 0, 3,
3, 0, 3, 1, 0, 1, 1, 3, 1, 3, 3, 0, 0, 1, 3, 0, 0, 3, 0, 1, 0, 1,
3, 1, 1, 0, 0, 3, 0, 1, 0, 0, 0, 1, 3, 0, 0, 1, 3, 3, 1, 3, 1, 3,
3, 3, 1, 3, 3, 1, 1, 1, 0, 0, 1, 1, 0, 3], dtype=int64)
```

```
In [52]: accuracy_score(y_test,y_pred)
```

```
Out[52]: 1.0
```

```
In [53]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	100
1	1.00	1.00	1.00	100
3	1.00	1.00	1.00	100
accuracy			1.00	300
macro avg	1.00	1.00	1.00	300
weighted avg	1.00	1.00	1.00	300

Step4. [Find Best no. of trees and Best Learning Rate using Grid Search and Cross Validation]

```
In [54]: Param_grid={'n_estimators':[50, 100, 200, 400],'learning_rate':[0.1, 0.01]}
```

```
In [55]: all_scores = cross_val_score(estimator=model,X=X_train,y=y_train,cv=5)
print(all_scores)
```

```
[0.99166667 1. 1. 1. 1.]
```

```
In [56]: model2 = GridSearchCV(estimator=model,param_grid=Param_grid,cv=5,n_jobs=-1)
```

```
In [57]: model2.fit(X_train,y_train)
```

```
Out[57]: GridSearchCV(cv=5,
                      estimator=GradientBoostingClassifier(learning_rate=1.0,
                                                            max_depth=1, max_features=4,
                                                            random_state=42,
                                                            subsample=0.5),
                      n_jobs=-1,
                      param_grid={'learning_rate': [0.1, 0.01],
                                   'n_estimators': [50, 100, 200, 400]})
```

```
In [58]: y_pred2=model2.predict(X_test)
         y_pred2
```

```
Out[58]: array([0, 1, 1, 1, 3, 1, 0, 0, 3, 3, 1, 3, 3, 3, 3, 1, 3, 3, 3, 0, 3, 1,
                0, 1, 0, 0, 0, 1, 0, 1, 3, 0, 3, 1, 1, 3, 0, 0, 3, 1, 0, 1, 0, 0,
                3, 0, 3, 3, 3, 1, 0, 0, 0, 1, 1, 0, 3, 3, 0, 3, 1, 1, 0, 1, 0, 3,
                3, 0, 3, 0, 0, 3, 0, 0, 1, 3, 3, 0, 3, 0, 3, 3, 1, 3, 1, 1, 1, 3,
                0, 3, 1, 3, 1, 1, 1, 3, 1, 0, 3, 1, 3, 0, 0, 3, 1, 3, 1, 0, 1, 0,
                1, 0, 0, 0, 1, 1, 1, 1, 3, 0, 0, 0, 0, 0, 1, 3, 3, 1, 1, 1, 1, 3,
                0, 1, 0, 3, 3, 3, 3, 3, 3, 1, 0, 3, 1, 1, 3, 3, 3, 0, 1, 1, 0, 0,
                0, 3, 0, 1, 3, 3, 1, 1, 1, 1, 0, 1, 0, 0, 0, 3, 0, 1, 1, 3, 0, 0,
                0, 0, 1, 3, 1, 0, 1, 1, 3, 1, 1, 0, 3, 0, 1, 3, 0, 3, 0, 1, 3, 0,
                1, 1, 3, 1, 0, 0, 0, 3, 3, 1, 1, 3, 3, 3, 3, 1, 0, 1, 0, 0, 3, 0,
                3, 3, 0, 0, 3, 1, 3, 0, 3, 1, 1, 3, 0, 1, 1, 1, 0, 0, 3, 1, 0, 3,
                3, 0, 3, 1, 0, 1, 1, 3, 1, 3, 3, 0, 0, 1, 3, 0, 0, 3, 0, 1, 0, 1,
                3, 1, 1, 0, 0, 3, 0, 1, 0, 0, 0, 1, 3, 0, 0, 1, 3, 3, 1, 3, 1, 3,
                3, 3, 1, 3, 3, 1, 1, 1, 0, 0, 1, 1, 0, 3], dtype=int64)
```

```
In [59]: accuracy_score(y_test,y_pred2)
```

```
Out[59]: 1.0
```

```
In [60]: print(classification_report(y_test,y_pred2))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	100
1	1.00	1.00	1.00	100
3	1.00	1.00	1.00	100
accuracy			1.00	300
macro avg	1.00	1.00	1.00	300
weighted avg	1.00	1.00	1.00	300

```
In [61]: model2.best_estimator_
```

```
Out[61]: GradientBoostingClassifier(max_depth=1, max_features=4, n_estimators=200,
                                     random_state=42, subsample=0.5)
```

Step5. [Build AdaBoostClassifier]

```
In [62]: base = DecisionTreeClassifier(max_features=4)
base2 = AdaBoostClassifier(base_estimator=base, random_state=0)
param_grid = {'n_estimators': [100, 150, 200], 'learning_rate': [0.01, 0.001]}
```

```
In [63]: model3 = GridSearchCV(base2, param_grid, cv=10, n_jobs=-1)
```

```
In [64]: model3.fit(X_train, y_train)
```

```
Out[64]: GridSearchCV(cv=10,
                      estimator=AdaBoostClassifier(base_estimator=DecisionTreeClassifier(
max_features=4),
                                                    random_state=0),
                      n_jobs=-1,
                      param_grid={'learning_rate': [0.01, 0.001],
'n_estimators': [100, 150, 200]})
```

```
In [65]: y_pred3=model3.predict(X_test)
y_pred3
```

```
Out[65]: array([0, 1, 1, 1, 3, 0, 0, 0, 3, 3, 1, 3, 3, 3, 3, 1, 3, 3, 3, 0, 3, 1,
0, 1, 0, 0, 0, 0, 0, 1, 3, 0, 3, 1, 1, 3, 0, 0, 3, 1, 1, 1, 0, 0,
3, 0, 3, 3, 3, 1, 1, 0, 0, 1, 1, 0, 3, 3, 1, 3, 1, 1, 0, 1, 0, 3,
3, 0, 3, 0, 0, 3, 0, 0, 1, 3, 3, 0, 3, 0, 3, 3, 1, 3, 1, 1, 1, 3,
0, 3, 1, 3, 1, 1, 1, 3, 1, 1, 3, 1, 3, 0, 0, 3, 0, 3, 1, 0, 1, 0,
1, 1, 0, 0, 0, 1, 1, 1, 3, 0, 0, 0, 0, 0, 1, 3, 3, 1, 1, 1, 1, 3,
1, 1, 0, 3, 3, 3, 3, 3, 3, 0, 0, 3, 1, 1, 3, 3, 3, 0, 1, 1, 0, 0,
0, 3, 0, 1, 3, 3, 1, 1, 0, 1, 0, 1, 1, 0, 0, 3, 0, 0, 1, 3, 0, 0,
1, 1, 1, 3, 0, 0, 1, 1, 3, 1, 1, 0, 3, 1, 1, 3, 0, 3, 1, 1, 3, 0,
1, 1, 3, 1, 1, 0, 0, 3, 3, 1, 1, 3, 3, 3, 3, 1, 0, 0, 0, 0, 3, 0,
3, 3, 0, 0, 3, 1, 3, 0, 3, 0, 1, 3, 0, 0, 1, 1, 0, 1, 3, 1, 0, 3,
3, 0, 3, 1, 0, 1, 1, 3, 1, 3, 3, 0, 0, 1, 3, 0, 0, 3, 0, 1, 0, 1,
3, 1, 0, 0, 0, 3, 0, 1, 0, 0, 1, 1, 3, 0, 0, 1, 3, 3, 1, 3, 1, 3,
3, 3, 1, 3, 3, 1, 1, 1, 0, 0, 1, 1, 0, 3], dtype=int64)
```

```
In [66]: accuracy_score(y_test, y_pred3)
```

```
Out[66]: 0.9133333333333333
```

```
In [67]: print(classification_report(y_test, y_pred3))
```

	precision	recall	f1-score	support
0	0.88	0.86	0.87	100
1	0.86	0.88	0.87	100
3	1.00	1.00	1.00	100
accuracy			0.91	300
macro avg	0.91	0.91	0.91	300
weighted avg	0.91	0.91	0.91	300

```
In [68]: model3.best_estimator_
```

```
Out[68]: AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_features=4),
                             learning_rate=0.01, n_estimators=100, random_state=0)
```

Step6. [Build LogisticRegressionCV classifier]

```
In [69]: model4 = LogisticRegressionCV(cv=4,Cs=5,penalty='l2')
```

```
In [70]: model4.fit(X_train,y_train)
```

```
Out[70]: LogisticRegressionCV(Cs=5, cv=4)
```

```
In [71]: y_pred4=model4.predict(X_test)
         y_pred4
```

```
Out[71]: array([0, 1, 1, 1, 3, 1, 0, 0, 3, 3, 1, 3, 3, 3, 3, 1, 3, 3, 3, 0, 3, 1,
                0, 1, 0, 0, 0, 1, 0, 1, 3, 0, 3, 1, 1, 3, 0, 0, 3, 1, 0, 1, 0, 0,
                3, 0, 3, 3, 3, 1, 0, 0, 0, 1, 1, 0, 3, 3, 0, 3, 1, 1, 0, 1, 0, 3,
                3, 0, 3, 0, 0, 3, 0, 0, 1, 3, 3, 0, 3, 0, 3, 3, 1, 3, 1, 1, 1, 3,
                0, 3, 1, 3, 1, 1, 1, 3, 1, 0, 3, 1, 3, 0, 0, 3, 1, 3, 1, 0, 1, 0,
                1, 0, 0, 0, 1, 1, 1, 1, 3, 0, 0, 0, 0, 0, 1, 3, 3, 1, 1, 1, 1, 3,
                0, 1, 0, 3, 3, 3, 3, 3, 3, 1, 0, 3, 1, 1, 3, 3, 3, 0, 1, 1, 0, 0,
                0, 3, 0, 1, 3, 3, 1, 1, 1, 1, 0, 1, 0, 0, 0, 3, 0, 1, 1, 3, 0, 0,
                0, 0, 1, 3, 1, 0, 1, 1, 3, 1, 1, 0, 3, 0, 1, 3, 0, 3, 0, 1, 3, 0,
                1, 1, 3, 1, 0, 0, 0, 3, 3, 1, 1, 3, 3, 3, 3, 1, 0, 1, 0, 0, 3, 0,
                3, 3, 0, 0, 3, 1, 3, 0, 3, 1, 1, 3, 0, 1, 1, 1, 0, 0, 3, 1, 0, 3,
                3, 0, 3, 1, 0, 1, 1, 3, 1, 3, 3, 0, 0, 1, 3, 0, 0, 3, 0, 1, 0, 1,
                3, 1, 1, 0, 0, 3, 0, 1, 0, 0, 0, 1, 3, 0, 0, 1, 3, 3, 1, 3, 1, 3,
                3, 3, 1, 3, 3, 1, 1, 1, 0, 0, 1, 1, 0, 3], dtype=int64)
```

```
In [72]: accuracy_score(y_test,y_pred4)
```

```
Out[72]: 1.0
```

```
In [73]: print(classification_report(y_test,y_pred4))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	100
1	1.00	1.00	1.00	100
3	1.00	1.00	1.00	100
accuracy			1.00	300
macro avg	1.00	1.00	1.00	300
weighted avg	1.00	1.00	1.00	300

Step 7 [Build VotingClassifier]

```
In [74]: model5=VotingClassifier(estimators=[('lr',model4),('gbc',base2)],voting='soft')
```

```
In [75]: model5.fit(X_train,y_train)
```

```
Out[75]: VotingClassifier(estimators=[('lr', LogisticRegressionCV(Cs=5, cv=4)),
                                      ('gbc',
                                       AdaBoostClassifier(base_estimator=DecisionTreeCla
ssifier(max_features=4),
                                                         random_state=0))],
                           voting='soft')
```

```
In [76]: y_pred5=model5.predict(X_test)
y_pred5
```

```
Out[76]: array([0, 1, 1, 1, 3, 0, 0, 0, 3, 3, 1, 3, 3, 3, 3, 1, 3, 3, 3, 0, 3, 1,
                0, 1, 0, 0, 0, 0, 0, 1, 3, 0, 3, 1, 1, 3, 0, 0, 3, 1, 1, 1, 0, 0,
                3, 0, 3, 3, 3, 1, 1, 0, 0, 1, 1, 0, 3, 3, 0, 3, 1, 1, 0, 1, 0, 3,
                3, 0, 3, 0, 0, 3, 0, 0, 1, 3, 3, 0, 3, 0, 3, 3, 1, 3, 1, 1, 1, 3,
                0, 3, 1, 3, 1, 1, 1, 3, 1, 0, 3, 1, 3, 0, 0, 3, 0, 3, 1, 0, 1, 0,
                1, 1, 0, 0, 0, 1, 1, 1, 3, 0, 0, 0, 0, 0, 1, 3, 3, 1, 1, 1, 1, 3,
                1, 1, 0, 3, 3, 3, 3, 3, 3, 0, 0, 3, 1, 1, 3, 3, 3, 0, 1, 1, 0, 0,
                0, 3, 0, 1, 3, 3, 1, 1, 0, 1, 0, 1, 0, 0, 0, 3, 0, 0, 1, 3, 0, 0,
                1, 1, 1, 3, 0, 0, 1, 1, 3, 1, 1, 0, 3, 1, 1, 3, 0, 3, 0, 1, 3, 0,
                1, 1, 3, 1, 0, 0, 0, 3, 3, 1, 1, 3, 3, 3, 3, 1, 0, 0, 0, 0, 3, 0,
                3, 3, 0, 0, 3, 1, 3, 0, 3, 0, 1, 3, 0, 0, 1, 1, 0, 1, 3, 1, 0, 3,
                3, 0, 3, 1, 0, 1, 1, 3, 1, 3, 3, 0, 0, 1, 3, 0, 0, 3, 0, 1, 0, 1,
                3, 1, 0, 0, 0, 3, 0, 1, 0, 0, 1, 1, 3, 0, 0, 1, 3, 3, 1, 3, 1, 3,
                3, 3, 1, 3, 3, 1, 1, 1, 0, 0, 1, 1, 0, 3], dtype=int64)
```

```
In [77]: accuracy_score(y_test,y_pred5)
```

```
Out[77]: 0.93
```

```
In [78]: print(classification_report(y_test,y_pred5))
```

	precision	recall	f1-score	support
0	0.88	0.91	0.90	100
1	0.91	0.88	0.89	100
3	1.00	1.00	1.00	100
accuracy			0.93	300
macro avg	0.93	0.93	0.93	300
weighted avg	0.93	0.93	0.93	300

Step8. [Interpret your results]

GradientBoostingClassifier(n_estimators=50)

GradientBoostingClassifier(n_estimators=50,learning_rate=1.0,max_depth=1,random_state=32)

In [79]: `model6 = GradientBoostingClassifier(n_estimators=50,learning_rate=1.0,max_depth=1,random_state=32)`

In [80]: `model6.fit(X_train,y_train)`

Out[80]: `GradientBoostingClassifier(learning_rate=1.0, max_depth=1, n_estimators=50, random_state=32)`

In [81]: `y_pred6=model6.predict(X_test)`
`y_pred6`

Out[81]: `array([0, 1, 1, 1, 3, 1, 0, 0, 3, 3, 1, 3, 3, 3, 3, 1, 3, 3, 3, 0, 3, 1, 0, 1, 0, 0, 3, 0, 3, 3, 3, 1, 0, 0, 0, 1, 1, 0, 3, 3, 0, 3, 1, 1, 0, 1, 0, 3, 3, 0, 3, 0, 0, 3, 0, 0, 1, 3, 3, 0, 3, 0, 3, 3, 1, 3, 1, 1, 1, 3, 0, 3, 1, 3, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 3, 0, 0, 0, 0, 0, 1, 3, 3, 1, 1, 1, 1, 1, 3, 0, 1, 0, 3, 3, 3, 3, 3, 3, 1, 0, 3, 1, 1, 3, 3, 3, 0, 1, 1, 0, 0, 0, 3, 0, 1, 3, 3, 1, 1, 1, 1, 0, 1, 0, 0, 0, 3, 0, 1, 1, 3, 0, 0, 0, 0, 1, 3, 1, 0, 1, 1, 3, 1, 1, 0, 3, 0, 1, 3, 0, 3, 0, 1, 3, 0, 1, 1, 3, 3, 3, 3, 1, 1, 3, 3, 3, 3, 1, 0, 1, 0, 0, 3, 0, 3, 3, 0, 0, 3, 1, 3, 0, 3, 1, 1, 3, 0, 1, 1, 1, 0, 0, 3, 1, 0, 3, 3, 0, 3, 1, 0, 1, 1, 3, 1, 3, 3, 0, 0, 1, 3, 0, 0, 3, 0, 1, 0, 1, 3, 1, 1, 0, 0, 3, 0, 1, 0, 0, 0, 1, 3, 0, 0, 1, 3, 3, 1, 3, 1, 3, 3, 1, 3, 3, 1, 1, 1, 0, 0, 1, 1, 0, 3], dtype=int64)`

In [82]: `accuracy_score(y_test,y_pred6)`

Out[82]: `1.0`

In [83]: `print(classification_report(y_test,y_pred6))`

	precision	recall	f1-score	support
0	1.00	1.00	1.00	100
1	1.00	1.00	1.00	100
3	1.00	1.00	1.00	100
accuracy			1.00	300
macro avg	1.00	1.00	1.00	300
weighted avg	1.00	1.00	1.00	300

AdaBoostClassifier

AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning_rate=0.01, n_estimators=75, random_state=0)

In [84]: `base3 = AdaBoostClassifier(base_estimator=base,learning_rate=0.01,n_estimators=75,random_state=0)`


```
In [85]: model7 = GridSearchCV(base3,param_grid,cv=5,n_jobs=-1)
```

```
In [86]: model7.fit(X_train,y_train)
```

```
Out[86]: GridSearchCV(cv=5,
                      estimator=AdaBoostClassifier(base_estimator=DecisionTreeClassifier
                      (max_features=4),
                      learning_rate=0.01, n_estimators=75,
                      random_state=0),
                      n_jobs=-1,
                      param_grid={'learning_rate': [0.01, 0.001],
                      'n_estimators': [100, 150, 200]})
```

```
In [87]: y_pred7=model7.predict(X_test)
y_pred7
```

```
Out[87]: array([0, 1, 1, 1, 3, 0, 0, 0, 3, 3, 1, 3, 3, 3, 3, 1, 3, 3, 3, 0, 3, 1,
                0, 1, 0, 0, 0, 0, 0, 1, 3, 0, 3, 1, 1, 3, 0, 0, 3, 1, 1, 1, 0, 0,
                3, 0, 3, 3, 3, 1, 1, 0, 0, 1, 1, 0, 3, 3, 1, 3, 1, 1, 0, 1, 0, 3,
                3, 0, 3, 0, 0, 3, 0, 0, 1, 3, 3, 0, 3, 0, 3, 3, 1, 3, 1, 1, 1, 3,
                0, 3, 1, 3, 1, 1, 1, 3, 1, 1, 3, 1, 3, 0, 0, 3, 0, 3, 1, 0, 1, 0,
                1, 1, 0, 0, 0, 1, 1, 1, 3, 0, 0, 0, 0, 0, 1, 3, 3, 1, 1, 1, 1, 3,
                1, 1, 0, 3, 3, 3, 3, 3, 3, 0, 0, 3, 1, 1, 3, 3, 3, 0, 1, 1, 0, 0,
                0, 3, 0, 1, 3, 3, 1, 1, 0, 1, 0, 1, 1, 0, 0, 3, 0, 0, 1, 3, 0, 0,
                1, 1, 1, 3, 0, 0, 1, 1, 3, 1, 1, 0, 3, 1, 1, 3, 0, 3, 1, 1, 3, 0,
                1, 1, 3, 1, 1, 0, 0, 3, 3, 1, 1, 3, 3, 3, 3, 1, 0, 0, 0, 0, 3, 0,
                3, 3, 0, 0, 3, 1, 3, 0, 3, 0, 1, 3, 0, 0, 1, 1, 0, 1, 3, 1, 0, 3,
                3, 0, 3, 1, 0, 1, 1, 3, 1, 3, 3, 0, 0, 1, 3, 0, 0, 3, 0, 1, 0, 1,
                3, 1, 0, 0, 0, 3, 0, 1, 0, 0, 1, 1, 3, 0, 0, 1, 3, 3, 1, 3, 1, 3,
                3, 3, 1, 3, 3, 1, 1, 1, 0, 0, 1, 1, 0, 3], dtype=int64)
```

```
In [88]: accuracy_score(y_test,y_pred7)
```

```
Out[88]: 0.9133333333333333
```

```
In [89]: print(classification_report(y_test,y_pred7))
```

	precision	recall	f1-score	support
0	0.88	0.86	0.87	100
1	0.86	0.88	0.87	100
3	1.00	1.00	1.00	100
accuracy			0.91	300
macro avg	0.91	0.91	0.91	300
weighted avg	0.91	0.91	0.91	300

