lab 5:-

step-1: import pandas as pd

import csv

```
dPab = pd.read_csv("diabetes.csv")
dPab.
dPab.head()
dPab.shape
df = pd.read_csv("diabetes.csv")
df.
y = df.columns
b df.info()
df.dtypes
df.dtypes.value_counts()
```

step2:-

```
X = df.drop("Outcome", axis=1)
Y = df[["Outcome"]]
X
Y
from sklearn.model_selection import.
    stratified_shuffle_split.
```

Lab5. Diabetes Classification using Logistic Regression

Objectives

In this lab, you will classify using Logistic Regression model, whether a person will become diabetic or not using PIMA diabetes dataset available in UCI. You will also predict a new person who is not in the dataset will become diabetic of not based on his details.

Learning Outcomes

After completing this lab, you will be able to

- Understand data and build baseline LogisticRegression model
- Create Heatmap with confusion matrix values
- Apply scaling using StandardScaler and MinMaxScaler and rebuild LoR model
- Print classification metrics scores and plot ROC curve
- Compare the performance of LoR model with LogisticRegressionCV with L1 and L2 regularization

Business Use Case

You are a data scientist. A leading hospital in your city has approached you with the medical details of their patients related to their diabetes information. The hospital has given you a file (diabetes.csv) that contains details of 768 patients and each patient is described with these 9 features. Here, ***Outcome*** is the ***dependent variable*** and ***all other 8 variables are independent variables.***

- Pregnancies: Number of times pregnant
- Glucose: Plasma glucose concentration over 2 hours in an oral glucose tolerance test
- BloodPressure: Diastolic blood pressure (mm Hg)
- SkinThickness: Triceps skin fold thickness (mm)
- Insulin: 2-Hour serum insulin (mu U/ml)
- BMI: Body mass index (weight in kg/(height in m)2)
- DiabetesPedigreeFunction: Diabetes pedigree function (a function which scores likelihood of diabetes based on family history)
- Age: Age (years)
- Outcome: Class variable (0 if non-diabetic, 1 if diabetic)

The hospital wants you to build a model so that the model will assess when a new patient visits them he will become diabetic or not.

Step1. [Understand Data]. Using Pandas, import **"diabetes.csv"** file and print properties such as head, shape, columns, dtype, info and value_counts.

Step2. [Build Logistic Regression Model]

- Prepare X matrix (8 feature columns) and y vector (ie., Outcome column)
- Split dataset with stratified shuffle split for training and testing as X_train, X_test, y_train, y_test (use **25%** test size).
- Create LogisticRegression model, fit on training set and predict on test set

Step3. [Predict on a new sample]

- Will this person become diabetic?. His details are given below.
- new_person = [[6, 200, 90, 10, 25, 23.3, 0.672, 42]]

Step3. [Compute Classification Metrics]

- Compute and print Accuracy, Precision, Recall and AUC scores

Step4. [Understand Correlation]

- Create confusion matrix between y_test and y_pred and plot confusion matrix values in a Heatmap. Explain the meaning of the 4 numbers you get.

stratified shuffle split()

```
shuf = StratifiedShuffleSplit(n_splits=4, test_size=0.25,
                              random_state=0)
shuf.get_n_splits(X,Y)
import warnings
warnings.filterwarnings("ignore")
for train, test in shuf.split(X,Y):
    X_train, X_test = X.iloc[train], X.iloc[test]
    y_train, y_test = y.iloc[train], y.iloc[test]
from sklearn.linear_model import LogisticRegression
    logmodel = LogisticRegression()
    logmodel.fit(X_train, Y_train)
      y_predic = logmodel.predict(X_test)
      y_predic
    logmodel.score(X_train, Y_train)
```

step 3:

```
new_person = [[6, 200, 90, 10, 25, 33.3, 0.672, 400]]
print(logmodel.predict(new_person))
```

Precision:

```
from sklearn.metrics import precision_score
print(precision_score(y_test, y_predic))
```

Recall:

```
from sklearn.metrics import recall_score
print(recall_score(y_test, y_predic))
```

Step5. [Normalization using MinmaxScaler and rebuild LoR]

- Now, normalize your X_train and X_test values using MinmaxScaler
- Create a new LogisticRegression model, fit on normalized training set and predict on the normalized test set
- Compute and print Accuracy, Precision, Recall and AUC scores

Step6. [Normalization using StandardScaler and rebuild LoR]

- Repeat Step5 with StandardScaler
- Among the 3 models, which model gives better classification scores?

Step7. [Plot ROC curve]

- Plot ROC curve as shown below. You can use the MinmaxScaler scaled values of X_test for computing predict_proba() score.

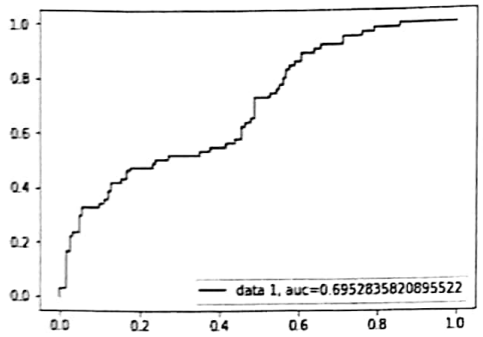

Step8. [Comparison with KNN classifier].

Create a KNN classifier with default values, fit on the scaled X using MinmaxScaler, predict and print classification metric scores.

Step9. [Update ROC curve]

- Update your ROC curve, this time, with one more curve of KNN classifier, as shown below.

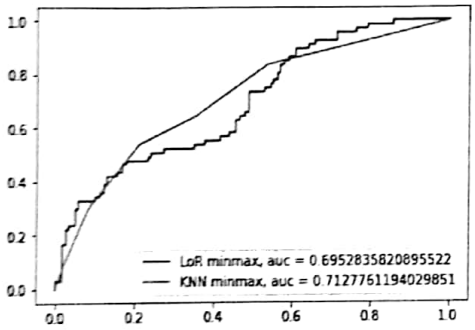

Step10. [Regularization]

- In order to reduce overfitting of your data, you will use *LogisticRegressionCV* model with L1 and L2 regularization parameters. Create both models using the following statements
  - `model1 = LogisticRegressionCV(Cs=10, cv=4, penalty='l1', solver='liblinear')`
  - `model2 = LogisticRegressionCV(Cs=10, cv=4, penalty='l2')`
- Perform fit using MinmaxScaler scaled values and predict

Step11. [Update ROC curve]

- Update your ROC curve, this time, with two more curves, as shown below

Accuracy:

```
from sklearn.metrices import accuracy_score
log_acscore = accuracy_score(y_test, y_predic)
log_acscore
```

AUC Scores

```
from sklearn.metrices import roc_auc_score
print(roc_auc_score(y_test, y_predic))
```

Step: 4

```
from sklearn.metrices import confusion_matrix
Confu_matrix = confusion_matrix(y_test, y_predic)
Confu_matrix
Confu_accu_score = accuracy_score(y_test, y_predic)
Confu_accu_score
import seaborn as sns
sns.heatmap(confusion_matrix(y_test, y_periodic)
/len(y), cmap="YlGnBu" annot=True)
```

Step: 5

```
from sklearn.preprocessing import minmax scalar
Scales = MinmaxScaler()
X_trained_min = Scaler.fit_transform(X_train)
X_test_min = Scaler.transform(x_test)
X_trained_min.shape
X_tested_min.shape
```

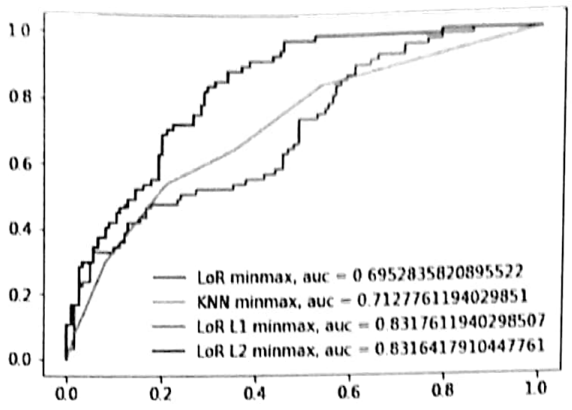

Out of these 4 models, which model performs the best?. How?. Why?.

```
from sklearn.linear_model import LogisticRegression.
logmodel1 = LogisticRegression()
logmodel1.fit(x_trained_min, y_train
y_predict = logmodel1.predict(x_tested_min)
y_predict.
logmodel1.score(x_trained.min, y_train)
```

Precision:

```
from Sklearn.metrices import precision_score
print(precision_score(y_test, y_predict)
```

Recall

```
from sklearn.metrices.import recall_score
print(recall_score(y_test, y_predict)
```

Accuracy.

```
from sklearn.metrices.import accuracy_score
min_accuscore = accuracy_score(y_test, y_predict)
min_arcScore
```

Auc Score:

```
from sklearn.metrics import roc_auc_score
log_auc_sc = roc_auc_score(y_test, y_predict)
log_auc1 = (LOR.minman, Auc =, log_acscore)
log_auc1.
```

step:6

```
from sklearn.preprocessing import standardscalex
scalex = standardscalex()
X_trained_stand = scalex.fit_transform(X_train)
X_tested_stand = scalex.transform(x_test)
X_trained_stand.shape
X_tested_stand.shape
from sklearn.linear_model import logistic Regression
logmodel2 = logistic Regression()
log model 2.fit(Xtrained_stand, y_train)
y_predict_stand = logmodel2.predict(X_tested_stand)
y_predict = stand
logmodel2.score(x_trained.stand, y_train)
```

precision score

```
from sklearn.metrics import precision_score.
print(precision_score(y_test, y_predict_stand))
```

NOTES

Recall score

```
from sklearn.metrices import recall_score
print (recall_score (y_test, y_predict_stand))
```

Accuracy:

```
from sklearn.metrices import accuracy_score
stand_accscore = accuracy_score(y_test, y_predict_stand)
stand_accuscore.
```

Auc score

```
from sklearn.metrices import roc_auc_score
stand_aucscore = roc_auc_score (y_test, y_predict_stand)
stand_aucs = (AUC= stand_aucscore)
stand aucs.
print ('logistic Regression model); log_accscore)
print ('Minmax Scales:', min_accuscore
print ('standard scaler:', stand_accscore)
```

Step 4:

```
from sklearn.metrices import roc_curve
pred_prb3 = log_model.predict_proba
(x_tested_min)
```

```
fpr3, tpr3, threshold = roc_curve
    (y_test, pred_prb3[:,1], pos_label=1)
import matplotlib.pyplot as plt
    plt.style.use('seaborn')
    plt.annotate(xy=[0.7,0], s=str(auc3))
plt.plot(fpr1, tpr1, linestyle=':', color='black',
                label='Logistic Regression')
plt.title('Roc curve')
plt.xlabel('False positive Rate')
plt.ylabel('True positive Rate')
```

step: 8:

```
from sklearn.neighbors import KNeighboursclassifier.
log_model3 = KNeighboursClassifier(n_neighbors=3)
log_model3.fit(x_trained_min, y_train)
    knn_y_pred = log_model3.predict(x_testedmin)
    knn_y_pred
```