

Name:Viviyan Richards W

#Roll no:205229133

Lab5. Diabetes Classification using Logistic Regression

Step1. [Understand Data]. Using Pandas, import “diabetes.csv” file and print properties such as head, shape, columns, dtype, info and value_counts

```
In [2]: import pandas as pd
data = pd.read_csv('diabetes.csv')
data.head()
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6			0.627
1	1	85	66	29	0	26.6			0.351
2	8	183	64	0	0	23.3			0.672
3	1	89	66	23	94	28.1			0.167
4	0	137	40	35	168	43.1			2.288

```
In [3]: data.shape
```

```
Out[3]: (768, 9)
```

```
In [4]: data.columns
```

```
Out[4]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [5]: data.dtypes
```

```
Out[5]: Pregnancies      int64
Glucose      int64
BloodPressure  int64
SkinThickness  int64
Insulin      int64
BMI          float64
DiabetesPedigreeFunction  float64
Age          int64
Outcome      int64
dtype: object
```

```
In [6]: type(data)
```

```
Out[6]: pandas.core.frame.DataFrame
```

```
In [7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [8]: data.value_counts()
```

```
Out[8]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesPedi
greeFunction  Age  Outcome
0      0      57      60      0      0      21.7  0.735
67  0      1      76      0      0      45.3  0.194
46  0      1      108      37      0      39.2  0.305
65  0      1      74      0      0      28.8  0.153
48  0      1      72      29      325      36.9  0.159
28  0      1
..
2      84      50      23      76      30.4  0.968
21  0      1      65      0      0      39.6  0.930
27  0      1      0      23      0      28.9  0.773
25  0      1      58      16      52      32.7  0.166
25  0      1      72      41      114      40.9  0.817
17      163      72
47  1      1
Length: 768, dtype: int64
```

Step2. [Build Logistic Regression Model]

Prepare X matrix (8 feature columns) and y vector (ie., Outcome column)

```
In [9]: X=data.drop("Outcome", axis=1)
        y=data[["Outcome"]]
```

```
In [10]: X
```

```
Out[10]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
...
763	10	101	76	48	180	32.9	0.171
764	2	122	70	27	0	36.8	0.340
765	5	121	72	23	112	26.2	0.245
766	1	126	60	0	0	30.1	0.349
767	1	93	70	31	0	30.4	0.315

768 rows × 8 columns



```
In [11]: y
```

```
Out[11]:
```

	Outcome
0	1
1	0
2	1
3	0
4	1
...	...
763	0
764	0
765	0
766	1
767	0

768 rows × 1 columns

Split dataset with stratified shuffle split for training and testing as X_{train} , X_{test} , y_{train} , y_{test} (use 25% test size).

```
In [12]: from sklearn.model_selection import StratifiedShuffleSplit
```

```
In [13]: StratifiedShuffleSplit()
shuffle = StratifiedShuffleSplit(n_splits=4, test_size=0.25, random_state=0)
```

```
In [14]: shuffle.get_n_splits(X,y)
```

```
Out[14]: 4
```

Create LogisticRegression model, fit on training set and predict on test set

```
In [15]: import warnings
warnings.filterwarnings('ignore')
```

```
In [16]: for train, test in shuffle.split(X, y):
          X_train, X_test = X.iloc[train], X.iloc[test]
          y_train, y_test = y.iloc[train], y.iloc[test]
```

```
In [17]: from sklearn.linear_model import LogisticRegression
```

```
In [18]: lmodel = LogisticRegression()
```

```
In [19]: lmodel.fit(X_train,y_train)
```

```
Out[19]: LogisticRegression()
```

```
In [20]: y_predict = lmodel.predict(X_test)
y_predict
```

```
Out[20]: array([0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
                1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
                1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
                0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
                1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
                0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
                0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0], dtype=int64)
```

```
In [21]: lmodel.score(X_train,y_train)
```

```
Out[21]: 0.7899305555555556
```

Predict on a new sample

```
In [22]: new_person = [[6, 200, 90, 10, 25, 23.3, 0.672, 42]]
```

```
In [23]: print(lmodel.predict(new_person))
```

```
[1]
```

Step3. [Compute Classification Metrics]

Precision score

```
In [24]: from sklearn.metrics import precision_score  
print(precision_score(y_test, y_predict))
```

```
0.6727272727272727
```

Recall score

```
In [25]: from sklearn.metrics import recall_score  
print(recall_score(y_test, y_predict))
```

```
0.5522388059701493
```

Accuracy score

Which is 75% Accurate

```
In [26]: from sklearn.metrics import accuracy_score  
lor_ascore=accuracy_score(y_test, y_predict)
```

```
In [27]: lor_ascore
```

```
Out[27]: 0.75
```

AUC score

AUC is the percentage of the ROC plot that is underneath the curve:

```
In [28]: from sklearn.metrics import roc_auc_score
```

```
In [29]: print(roc_auc_score(y_test, y_predict))
```

```
0.7041194029850747
```

Step4. [Understand Correlation]

Create confusion matrix between y_{test} and y_{pred} and plot confusion matrix values in a Heatmap. Explain the meaning of the 4 numbers you get.

```
In [30]: from sklearn.metrics import confusion_matrix
cnf_matrix=confusion_matrix(y_test, y_predict)
cnf_matrix
```

```
Out[30]: array([[107, 18],
               [ 30, 37]], dtype=int64)
```


```
In [31]: cf_ac_score = accuracy_score(y_test, y_predict)
```

```
In [32]: cf_ac_score
```

```
Out[32]: 0.75
```

```
In [33]: import seaborn as sns
sns.heatmap(confusion_matrix(y_test,y_predict) / len(y), cmap='YlGnBu', annot=True)
```

```
Out[33]: <AxesSubplot:>
```



X-axis is predicted value

Y-axis is real value

Step5. [Normalization using MinmaxScaler and rebuild LoR]

Now, normalize your X_{train} and X_{test} values using MinmaxScaler

```
In [34]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_trained = scaler.fit_transform(X_train)
X_tested = scaler.transform(X_test)
```

```
In [35]: X_trained.shape
```

```
Out[35]: (576, 8)
```

```
In [36]: X_tested.shape
```

```
Out[36]: (192, 8)
```

Create a new LogisticRegression model, fit on normalized training set and predict on the normalized test set

```
In [37]: from sklearn.linear_model import LogisticRegression  
lmodel1 = LogisticRegression()
```

```
In [38]: lmodel1.fit(X_trained, y_train)
```

```
Out[38]: LogisticRegression()
```

```
In [39]: yn_predict = lmodel1.predict(X_tested)  
yn_predict
```

```
Out[39]: array([0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,  
                0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,  
                1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,  
                1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,  
                0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,  
                1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,  
                0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,  
                0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0]) dtype=int64)
```

```
In [40]: lmodel1.score(X_trained, y_train)
```

```
Out[40]: 0.7899305555555556
```

Precision score

```
In [41]: from sklearn.metrics import precision_score  
print(precision_score(y_test, yn_predict))
```

```
0.673469387755102
```

Recall score

```
In [42]: from sklearn.metrics import recall_score  
print(recall_score(y_test, yn_predict))
```

```
0.4925373134328358
```

Accuracy score

Which is 72% Accurate

```
In [43]: from sklearn.metrics import accuracy_score  
minmax_ascore=accuracy_score(y_test, yn_predict)
```

```
In [44]: minmax_ascore
```

```
Out[44]: 0.7395833333333334
```

AUC score

AUC is the percentage of the ROC plot that is underneath the curve

```
In [45]: from sklearn.metrics import roc_auc_score
```

```
In [46]: lgr_auc=roc_auc_score(y_test, yn_predict)  
lgr_auc1=('LoR minmax, AUC=',lgr_auc)  
lgr_auc1
```

```
Out[46]: ('LoR minmax, AUC=', 0.6822686567164179)
```

Step6. [Normalization using StandardScaler and rebuild LoR]

```
In [47]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
Xs_trained = scaler.fit_transform(X_train)  
Xs_tested = scaler.transform(X_test)
```

```
In [48]: Xs_trained.shape
```

```
Out[48]: (576, 8)
```

```
In [49]: Xs_tested.shape
```

```
Out[49]: (192, 8)
```

```
In [50]: from sklearn.linear_model import LogisticRegression  
lmodel2 = LogisticRegression()
```

```
In [51]: lmodel2.fit(Xs_trained, y_train)
```

```
Out[51]: LogisticRegression()
```



```
In [52]: ys_predict = lmodel2.predict(Xs_tested)
ys_predict
```

```
Out[52]: array([0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
                1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
                1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
                0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
                1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0,
                0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0,
                0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0], dtype=int64)
```

```
In [53]: lmodel2.score(Xs_train, y_train)
```

```
Out[53]: 0.7795138888888888
```

Precision score

```
In [54]: from sklearn.metrics import precision_score
print(precision_score(y_test, ys_predict))
```

```
0.6851851851851852
```

Recall score

```
In [55]: from sklearn.metrics import recall_score
print(recall_score(y_test, ys_predict))
```

```
0.5522388059701493
```

Accuracy score

Which is 75% Accurate

```
In [56]: from sklearn.metrics import accuracy_score
ss_ascore=accuracy_score(y_test, ys_predict)
```

```
In [57]: ss_ascore
```

```
Out[57]: 0.7552083333333334
```

AUC score

AUC is the percentage of the ROC plot that is underneath the curve

```
In [58]: from sklearn.metrics import roc_auc_score
```

```
In [59]: ss_auc=roc_auc_score(y_test, ys_predict)
ss_auc1=('AUC=',ss_auc)
ss_auc1
```

```
Out[59]: ('AUC=', 0.7081194029850746)
```

Among the 3 models, which model gives better classification scores?

```
In [60]: print('StandardScaler:',ss_ascore)
print('MinmaxScaler:',minmax_ascore)
print('Logistic Regression Model:',lor_ascore)
```

```
StandardScaler: 0.7552083333333334
```

```
MinmaxScaler: 0.7395833333333334
```

```
Logistic Regression Model: 0.75
```

Step7. [Plot ROC curve]


Plot ROC curve as shown below. You can use the MinmaxScaler scaled values of X_test for computing predict_proba() score.

```
In [61]: from sklearn.metrics import roc_curve
```

```
In [62]: predict_pb1 = lmodel1.predict_proba(X_tested)
fpr1, tpr1, threshold1 = roc_curve(y_test,predict_pb1[:,1], pos_label=1)
```

```
In [63]: import matplotlib.pyplot as plt
plt.style.use('seaborn')
plt.annotate(xy=[0.7,0], s=ss_auc1)
plt.plot(fpr1, tpr1, linestyle=':', color='black',label='Logistic Regression')
plt.title('Roc curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

Out[63]: Text(0, 0.5, 'True Positive Rate')



Step8. [Comparison with KNN classifier].

Create a KNN classifier with default values, fit on the scaled X using MinmaxScaler, predict and print classification metric scores.

```
In [64]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [65]: lmodel3 = KNeighborsClassifier(n_neighbors=3)
```

```
In [66]: lmodel3.fit(X_train,y_train)
```

Out[66]: KNeighborsClassifier(n_neighbors=3)

```
In [67]: knn_y_predict = lmodel3.predict(X_tested)
knn_y_predict
```

```
Out[67]: array([1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
        0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
        1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
        0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0,
        1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0,
        0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0,
        0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1], dtype=int64)
```

Precision score

```
In [68]: from sklearn.metrics import precision_score
print(precision_score(y_test, knn_y_predict))
```

```
0.6271186440677966
```

Recall score

```
In [69]: from sklearn.metrics import recall_score
print(recall_score(y_test, knn_y_predict))
```

```
0.5522388059701493
```

Accuracy score

```
In [70]: from sklearn.metrics import accuracy_score
kn_ascore=accuracy_score(y_test, knn_y_predict)
kn_ascore
```

```
Out[70]: 0.7291666666666666
```

AUC score

```
In [71]: from sklearn.metrics import roc_auc_score
kn_auc=roc_auc_score(y_test, knn_y_predict)
kn_auc1=('KNN minmax, AUC=',kn_auc)
kn_auc1
```

```
Out[71]: ('KNN minmax, AUC=', 0.6881194029850747)
```


Step9. [Update ROC curve]

Update your ROC curve, this time, with one more curve of KNN classifier,

```
In [72]: predict_pb2 = lmodel3.predict_proba(X_tested)
fpr1,tpr1,threshold1 = roc_curve(y_test, predict_pb1[:,1],pos_label=1)
fpr2,tpr2,threshold2 = roc_curve(y_test, predict_pb2[:,1],pos_label=1)
```

```
In [73]: plt.plot(fpr1,tpr1,linestyle=':',color='black',label='Logistic Regression')
plt.plot(fpr2,tpr2,linestyle=':',color='blue',label='KNN')
plt.annotate(xy=[0.6,0.1], s=kn_auc1)
plt.annotate(xy=[0.6,0], s=lgr_auc1)
plt.legend(loc='best')
plt.title('Roc curve')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
```

Out[73]: Text(0, 0.5, 'True positive rate')



Step10. [Regularization]

In order to reduce overfitting of your data, you will use LogisticRegressionCV model with L1 and L2 regularization parameters. Create both models using the following statements

```
In [74]: from sklearn.linear_model import LogisticRegressionCV
```

```
In [75]: lmodel4 = LogisticRegressionCV(Cs=10, cv=4, penalty='l1', solver='liblinear')
lmodel5 = LogisticRegressionCV(Cs=10, cv=4, penalty='l2')
```

Perform fit using MinmaxScaler scaled values and predict

```
In [76]: print(lmodel4.fit(X_train,y_train))
print(lmodel5.fit(X_train,y_train))
```

```
LogisticRegressionCV(cv=4, penalty='l1', solver='liblinear')
LogisticRegressionCV(cv=4)
```

```
In [77]: lr_y_predict1 = lmodel4.predict(X_tested)
lr_y_predict2 = lmodel5.predict(X_tested)
print('Logistic RegressionCV L1:\n',lr_y_predict1)
print('-----')
print('Logistic RegressionCV L2:\n',lr_y_predict2)
```

```
Logistic RegressionCV L1:
```

```
[0 0 1 1 0 1 1 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 1 0 0 0 0 0 0
0 0 0 0 1 0 1 1 0 1 0 1 0 0 0 1 0 0 1 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0
1 0 0 0 0 1 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 0
0 0 1 1 0 0 0 0 0 1 0 0 1 1 0 0 0 0 0 0 0 1 0 1 1 1 1 0 0 0 0 0 0 1 0 0 0
1 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 0 0 0 0 0 1 0 0 1 1 0 0
0 0 0 1 0 0 0]
```

```
-----
Logistic RegressionCV L2:
```

```
[0 0 1 1 0 1 1 0 0 1 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 1 0 0 0 0 0 0
0 0 0 0 1 0 1 1 0 1 0 1 0 0 0 1 0 0 1 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 1
1 0 0 0 0 1 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 0
0 0 1 1 0 0 0 0 0 1 0 0 1 1 0 0 0 0 0 0 0 1 0 1 1 1 1 0 0 0 1 0 0 1 0 0 0
1 1 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0 0 0 0 1 0 0 1 1 0 0
0 0 0 1 0 0 0]
```

```
In [78]: from sklearn.metrics import roc_auc_score
lrp_auc=roc_auc_score(y_test, lr_y_predict1)
lrp1_auc=('LoR L1 minmax, AUC=',lrp_auc)
lrp1_auc
```

```
Out[78]: ('LoR L1 minmax, AUC=', 0.6748059701492537)
```

```
In [79]: from sklearn.metrics import roc_auc_score
lrp_auc2=roc_auc_score(y_test, lr_y_predict2)
lrp2_auc=('LoR L2 minmax, AUC=',lrp_auc2)
lrp2_auc
```

```
Out[79]: ('LoR L2 minmax, AUC=', 0.6931940298507463)
```

Step11. [Update ROC curve]

Update your ROC curve, this time, with two more curves

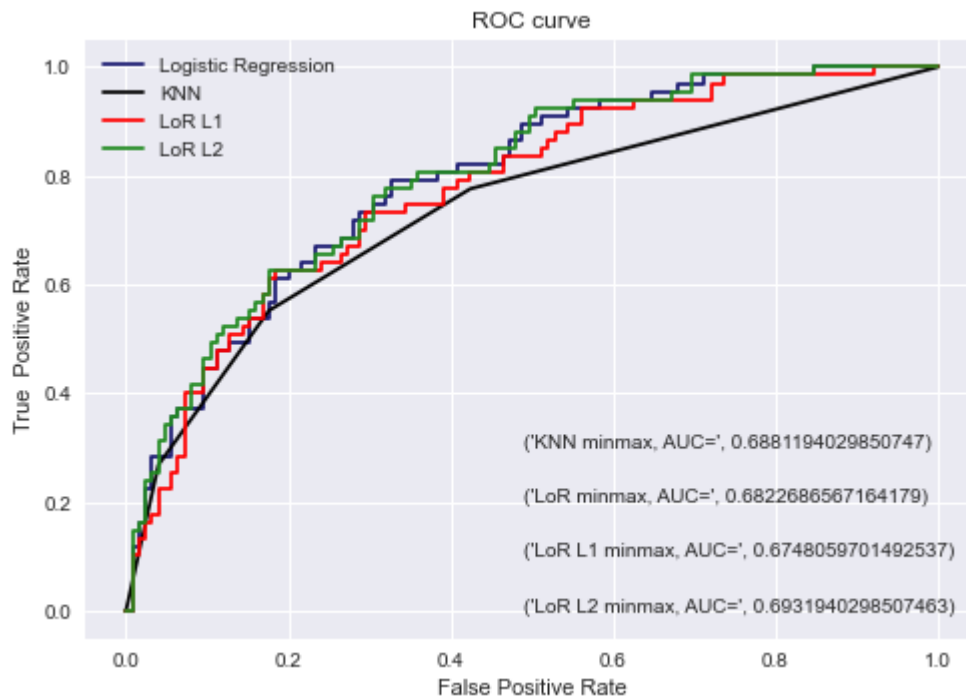
```
In [80]: predict_pb3 = lmodel4.predict_proba(X_tested)
predict_pb4 = lmodel5.predict_proba(X_tested)
fpr1,tpr1,thresh1 = roc_curve(y_test, predict_pb1[:,1], pos_label=1)
fpr2,tpr2,thresh2 = roc_curve(y_test, predict_pb2[:,1], pos_label=1)
fpr3,tpr3,thresh3 = roc_curve(y_test, predict_pb3[:,1], pos_label=1)
fpr4,tpr4,thresh4 = roc_curve(y_test, predict_pb4[:,1], pos_label=1)
```

```

In [90]: plt.plot(fpr1,tpr1,linestyle='-',color='midnightblue', label='Logistic Regression')
plt.plot(fpr2,tpr2,linestyle='-',color='black', label='KNN')
plt.plot(fpr3,tpr3,linestyle='-',color='red', label='LoR L1')
plt.plot(fpr4,tpr4,linestyle='-',color='forestgreen', label='LoR L2')
plt.annotate(xy=[0.49,0.3], s=kn_auc1)
plt.annotate(xy=[0.49,0.2], s=lgr_auc1)
plt.annotate(xy=[0.49,0.1], s=lrp1_auc)
plt.annotate(xy=[0.49,0], s=lrp2_auc)
plt.legend(loc='best')
plt.title('ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

```

Out[90]: Text(0, 0.5, 'True Positive Rate')



Logistic Regression L2 Minmax gives 69% of AUC score so this shows it is the best score