

```
In [133]: import pandas as pd  
import csv
```

```
In [134]: fuel = pd.read_csv('fuel_data.csv')  
fuel
```

Out[134]:

	drivenKM	fuelAmount
0	390.00	3600.0
1	403.00	3705.0
2	396.50	3471.0
3	383.50	3250.5
4	321.10	3263.7
5	391.30	3445.2
6	386.10	3679.0
7	371.80	3744.5
8	404.30	3809.0
9	392.20	3905.0
10	386.43	3874.0
11	395.20	3910.0
12	381.00	4020.7
13	372.00	3622.0
14	397.00	3450.5
15	407.00	4179.0
16	372.40	3454.2
17	375.60	3883.8
18	399.00	4235.9

```
In [135]: fuel.head()
```

Out[135]:

	drivenKM	fuelAmount
0	390.0	3600.0
1	403.0	3705.0
2	396.5	3471.0
3	383.5	3250.5
4	321.1	3263.7

```
In [136]: fuel.shape
```

Out[136]: (19, 2)

```
In [137]: fuel.columns
```

```
Out[137]: Index(['drivenKM', 'fuelAmount'], dtype='object')
```

```
In [138]: fuel.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19 entries, 0 to 18
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   drivenKM    19 non-null    float64
 1   fuelAmount  19 non-null    float64
dtypes: float64(2)
memory usage: 432.0 bytes
```

```
In [139]: fuel.isnull()
```


```
Out[139]:
```

	drivenKM	fuelAmount
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
5	False	False
6	False	False
7	False	False
8	False	False
9	False	False
10	False	False
11	False	False
12	False	False
13	False	False
14	False	False
15	False	False
16	False	False
17	False	False
18	False	False

```
In [140]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [141]: sns.relplot(x="drivenKM", y="fuelAmount", data=fuel)
```

```
Out[141]: <seaborn.axisgrid.FacetGrid at 0x23b5cb31550>
```



```
In [142]: data1 = ['drivenKM']  
X=fuel[data1]
```

```
In [143]: data2 = ['fuelAmount']  
Y=fuel[data2]
```

```
In [144]: print(X)  
X.dtypes
```

```
      drivenKM  
0      390.00  
1      403.00  
2      396.50  
3      383.50  
4      321.10  
5      391.30  
6      386.10  
7      371.80  
8      404.30  
9      392.20  
10     386.43  
11     395.20  
12     381.00  
13     372.00  
14     397.00  
15     407.00  
16     372.40  
17     375.60  
18     399.00
```

```
Out[144]: drivenKM      float64  
dtype: object
```

```
In [145]: print(Y)  
Y.dtypes  
y_True=y
```

```
      fuelAmount  
0      3600.0  
1      3705.0  
2      3471.0  
3      3250.5  
4      3263.7  
5      3445.2  
6      3679.0  
7      3744.5  
8      3809.0  
9      3905.0  
10     3874.0  
11     3910.0  
12     4020.7  
13     3622.0  
14     3450.5  
15     4179.0  
16     3454.2  
17     3883.8  
18     4235.9
```

```
In [146]: from sklearn.model_selection import train_test_split
```

```
In [147]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y,train_size=0.8,test_size=
```

```
In [148]: X_train, X_test, Y_train, Y_test
```

```
Out[148]: (      drivenKM
          9      392.20
          8      404.30
          5      391.30
          1      403.00
         14      397.00
          4      321.10
         16      372.40
          6      386.10
          7      371.80
         15      407.00
         10      386.43
         13      372.00
         17      375.60
         18      399.00
          3      383.50,
          drivenKM
         11      395.2
         12      381.0
          0      390.0
          2      396.5,
          fuelAmount
          9      3905.0
          8      3809.0
          5      3445.2
          1      3705.0
         14      3450.5
          4      3263.7
         16      3454.2
          6      3679.0
          7      3744.5
         15      4179.0
         10      3874.0
         13      3622.0
         17      3883.8
         18      4235.9
          3      3250.5,
          fuelAmount
         11      3910.0
         12      4020.7
          0      3600.0
          2      3471.0)
```

```
In [149]: X_train.shape
```

```
Out[149]: (15, 1)
```

```
In [150]: X_test.shape
```

```
Out[150]: (4, 1)
```

```
In [151]: Y_train.shape
```

```
Out[151]: (15, 1)
```

```
In [152]: Y_test.shape
```

```
Out[152]: (4, 1)
```

```
In [153]: from sklearn.linear_model import LinearRegression
```

```
In [154]: lin_reg = LinearRegression()  
lin_reg.fit(X_train,Y_train)
```

```
Out[154]: LinearRegression()
```

```
In [155]: x=[[800]]  
lin_reg.predict(x)
```

```
Out[155]: array([[6857.09812249]])
```

```
In [156]: y_pred=lin_reg.predict(X)  
y_pred
```

```
Out[156]: array([[3744.25861846],  
                 [3842.95840761],  
                 [3793.60851303],  
                 [3694.90872388],  
                 [3221.14973595],  
                 [3754.12859737],  
                 [3714.64868171],  
                 [3606.07891365],  
                 [3852.82838653],  
                 [3760.9616597 ],  
                 [3717.1541379 ],  
                 [3783.73853412],  
                 [3675.9279952 ],  
                 [3607.59737194],  
                 [3797.40465877],  
                 [3873.3275735 ],  
                 [3610.63428853],  
                 [3634.92962124],  
                 [3812.58924172]])
```

```
In [157]: from sklearn.metrics import mean_squared_error, r2_score
```

```
mse = mean_squared_error(y_True, y_pred)  
mse
```

```
Out[157]: 58987.44292405615
```

```
In [158]: r2_score(y_True, y_pred)
```

```
Out[158]: 0.21643923225776895
```

```
In [159]: lin_reg.coef_
```

```
Out[159]: array([[7.59229147]])
```

```
In [160]: lin_reg.intercept_
```

```
Out[160]: array([783.2649439])
```

```
In [161]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
s_X_train=scaler.fit_transform(X_train)  
s_X_train
```

```
Out[161]: array([[ 0.39182783],  
 [ 0.98313698],  
 [ 0.34784616],  
 [ 0.91960789],  
 [ 0.62639675],  
 [-3.08272426],  
 [-0.57576896],  
 [ 0.09372983],  
 [-0.60509007],  
 [ 1.11508199],  
 [ 0.10985644],  
 [-0.59531637],  
 [-0.41938968],  
 [ 0.7241338 ],  
 [-0.03332833]])
```

```
In [162]: s_X_test=scaler.transform(X_test)  
s_X_test
```

```
Out[162]: array([[ 0.5384334 ],  
 [-0.15549964],  
 [ 0.28431708],  
 [ 0.60196248]])
```

```
In [163]: s_lin_reg = LinearRegression()  
s_lin_reg.fit( s_X_train,y_train)
```

```
Out[163]: LinearRegression()
```

```
In [164]: s_y_pred=s_lin_reg.predict(s_X_test)  
s_y_pred
```

```
Out[164]: array([[3745.68215321],  
 [3676.96914641],  
 [3720.51964368],  
 [3751.97278059]])
```

```
In [165]: mean_squared_error(y_test, s_y_pred)
```

```
Out[165]: 67297.78117999973
```

```
In [166]: r2_score(y_test, s_y_pred)
```

```
Out[166]: 0.0345055936007026
```

```
In [167]: s_lin_reg.coef_
```

```
Out[167]: array([[99.01964899]])
```


```
In [168]: s_lin_reg.intercept_
```

```
Out[168]: array([3692.36666667])
```

```
In [169]: import matplotlib.pyplot as plt
import numpy as np

plt.scatter(y_True,y_pred)
```

```
Out[169]: <matplotlib.collections.PathCollection at 0x23b5cf242b0>
```



```
In [170]: from sklearn.preprocessing import MinMaxScaler
scalers = MinMaxScaler()
m_X_train=scalers.fit_transform(X_train)
m_X_train
```

```
Out[170]: array([[0.82770664],
 [0.9685681 ],
 [0.81722934],
 [0.95343423],
 [0.88358556],
 [0.        ],
 [0.59720605],
 [0.75669383],
 [0.59022119],
 [1.        ],
 [0.76053551],
 [0.59254948],
 [0.63445867],
 [0.90686845],
 [0.72642608]])
```



```
In [171]: m_X_test=scalers.transform(X_test)
m_X_test
```

```
Out[171]: array([[0.86263097],
                 [0.69732247],
                 [0.80209546],
                 [0.87776484]])
```

```
In [172]: m_lin_reg = LinearRegression()
m_lin_reg.fit( m_X_train,y_train)
```

```
Out[172]: LinearRegression()
```

```
In [173]: m_y_pred=m_lin_reg.predict(m_X_test)
m_y_pred
```

```
Out[173]: array([[3745.68215321],
                 [3676.96914641],
                 [3720.51964368],
                 [3751.97278059]])
```

```
In [174]: mean_squared_error(y_test, m_y_pred)
```

```
Out[174]: 67297.7811799998
```

```
In [175]: r2_score(y_test, m_y_pred)
```

```
Out[175]: 0.03450559360070149
```

```
In [176]: import numpy as np
from sklearn.neighbors import KNeighborsRegressor
neigh = KNeighborsRegressor(n_neighbors=5)
neigh.fit(X, y)
```

```
Out[176]: KNeighborsRegressor()
```

```
In [177]: n_y_pred=neigh.predict(X)
n_y_pred
```

```
Out[177]: array([[3700.64],
 [3875.88],
 [3794.48],
 [3684.84],
 [3593.64],
 [3746.84],
 [3684.84],
 [3745.04],
 [3875.88],
 [3666.24],
 [3569.74],
 [3794.48],
 [3741.6 ],
 [3745.04],
 [3794.48],
 [3875.88],
 [3745.04],
 [3745.04],
 [3754.48]])
```

```
In [178]: knn_mse=mean_squared_error(y_True, n_y_pred)
knn_mse
```

```
Out[178]: 70460.30507368421
```

```
In [179]: r2_score(y_True,n_y_pred)
```

```
Out[179]: 0.06403925984775638
```

```
In [180]: import numpy as np
from sklearn.linear_model import SGDRegressor
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
# Always scale the input. The most convenient way is to use a pipeline.
reg = make_pipeline(StandardScaler(), SGDRegressor(max_iter=1000, tol=1e-3))
reg.fit(X, y)
```

C:\Users\VISSWES\anaconda3\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(**kwargs)
```

```
Out[180]: Pipeline(steps=[('standardscaler', StandardScaler()),
 ('sgdregressor', SGDRegressor())])
```

```
In [181]: r_y_pred=reg.predict(X)
r_y_pred
```

```
Out[181]: array([3740.41973845, 3830.0665893 , 3785.24316387, 3695.59631303,
                3265.29142895, 3749.38442353, 3713.52568319, 3614.91414726,
                3839.03127438, 3755.59074398, 3715.80133402, 3776.27847879,
                3678.35653402, 3616.29332958, 3788.69111968, 3857.65023571,
                3619.05169422, 3641.11861136, 3802.48294288])
```

```
In [182]: sgd_mse=mean_squared_error(y_True,r_y_pred)
sgd_mse
```

```
Out[182]: 58824.133120863604
```

```
In [183]: r2_score(y_True,r_y_pred)
```

```
Out[183]: 0.218608560989888
```

```
In [184]: print('LR:',mse)
print('KNNR:',knn_mse)
print('SGDR:',sgd_mse)
```

```
LR: 58987.44292405615
KNNR: 70460.30507368421
SGDR: 58824.133120863604
```

```
In [ ]:
```

```
In [ ]:
```